

SARAH CONNOR

AI Security Engineer · Adversarial ML & Detection

✉ sarah@example.com · 📍 Los Angeles, CA

🌐 example.com · 📄 github.com/sconnor · 🔗 linkedin.com/in/sconnor

> SUMMARY

AI security engineer specializing in adversarial ML detection and red-teaming for production LLM systems. Shipped detection infrastructure handling 50k+ req/s with sub-second alert latency.

> EXPERIENCE



Senior AI Security Engineer

Cyberdyne Systems

Jun 2023 – Present

Los Angeles, CA

Cut adversarial-example detection latency 4x by rewriting the [batch, heads, seq, seq] scoring path to run inline with inference rather than as a post-hoc batch job.

- Built a red-team harness that found 12 prompt-injection classes missed by the prior eval suite, raising the pre-release catch rate from 61% to 94%.

Custom mutation-based fuzzer seeded from a 12k-prompt corpus, wired into the pre-release CI gate; findings triaged weekly with the app-sec team.

- Led migration of the model-serving detection pipeline to a streaming architecture, cutting p99 alert latency from 4.2s to 380ms.
- Mentored two junior engineers on adversarial evaluation methodology; both now own detection surfaces independently.

PyTorch · vLLM · Triton · Kubernetes



AI Security Engineer

Widget Robotics

Jan 2021 – May 2023

Remote

Owned anomaly scoring for a fleet-telemetry pipeline processing 50k events/s.

- Shipped an anomaly scoring service that cut false-positive alert volume 70% by replacing static thresholds with a learned per-device baseline.
- Built the eval harness now used across three teams to benchmark detection recall against a held-out attack corpus before every model release.

Python · scikit-learn · Kafka · Terraform



Machine Learning Engineer

Nova Systems Analytics

Aug 2019 – Dec 2020

San Diego, CA

- Trained and deployed a fraud-scoring model that reduced chargeback losses 18% in its first quarter in production.
- Instrumented model-serving latency tracing, cutting p99 debugging time from days to hours for the on-call rotation.

> SKILLS

Detection	Adversarial example detection, Anomaly scoring, Red-teaming, Eval design
Adversarial ML	Prompt injection, Model extraction, Data poisoning
Infra	Kubernetes, Terraform, Triton, CUDA, Kafka

> EDUCATION



MS Computer Science

University of California

2019

San Diego, US



CERTIFICATIONS



OSCP

Offensive Security Certified Professional

2022