

# SARAH CONNOR

AI Security Engineer · Adversarial ML & Detection  
sarah@example.com · Los Angeles, CA  
example.com · github.com/sconnor · linkedin.com/in/sconnor

## SUMMARY

AI security engineer specializing in adversarial ML detection and red-teaming for production LLM systems. Shipped detection infrastructure handling 50k+ req/s with sub-second alert latency.

## EXPERIENCE

**Senior AI Security Engineer** Jun 2023 – Present  
*Cyberdyne Systems* Los Angeles, CA

Cut adversarial-example detection latency 4x by rewriting the [batch, heads, seq, seq] scoring path to run inline with inference rather than as a post-hoc batch job.

- Built a red-team harness that found 12 prompt-injection classes missed by the prior eval suite, raising the pre-release catch rate from 61% to 94%.  
Custom mutation-based fuzzer seeded from a 12k-prompt corpus, wired into the pre-release CI gate; findings triaged weekly with the app-sec team.
  - Led migration of the model-serving detection pipeline to a streaming architecture, cutting p99 alert latency from 4.2s to 380ms.
  - Mentored two junior engineers on adversarial evaluation methodology; both now own detection surfaces independently.
- PyTorch · vLLM · Triton · Kubernetes

**AI Security Engineer** Jan 2021 – May 2023  
*Widget Robotics* Remote

Owned anomaly scoring for a fleet-telemetry pipeline processing 50k events/s.

- Shipped an anomaly scoring service that cut false-positive alert volume 70% by replacing static thresholds with a learned per-device baseline.
  - Built the eval harness now used across three teams to benchmark detection recall against a held-out attack corpus before every model release.
- Python · scikit-learn · Kafka · Terraform

**Machine Learning Engineer** Aug 2019 – Dec 2020  
*Nova Systems Analytics* San Diego, CA

- Trained and deployed a fraud-scoring model that reduced chargeback losses 18% in its first quarter in production.
- Instrumented model-serving latency tracing, cutting p99 debugging time from days to hours for the on-call rotation.

## SKILLS

|                |                                                                          |
|----------------|--------------------------------------------------------------------------|
| Detection      | Adversarial example detection, Anomaly scoring, Red-teaming, Eval design |
| Adversarial ML | Prompt injection, Model extraction, Data poisoning                       |
| Infra          | Kubernetes, Terraform, Triton, CUDA, Kafka                               |

## EDUCATION

**MS Computer Science** 2019  
*University of California* San Diego, US

## CERTIFICATIONS

---

### OSCP

*Offensive Security Certified Professional*

2022