

SARAH CONNOR

AI Security Engineer · Adversarial ML & Detection

sarah@example.com · Los Angeles, CA

example.com · github.com/sconnor

> 28 August 2026

Dear Hiring Manager,

I have spent the last three years building detection systems that assume the adversary reads the same papers I do. Vantage Security's recent work on prompt-injection taxonomies is the first public research I've seen that treats this as an adversarial classification problem rather than a filtering problem, and I would like to help build it.

At Cyberdyne Systems, I led the migration of our model-serving detection pipeline from a nightly batch job to a streaming architecture — a decision that looked obvious in retrospect and was not obvious at the time. The batch job had grown 40% slower every quarter as traffic scaled, and the team's first instinct was to throw more compute at it. I spent two weeks profiling instead, found the actual bottleneck was a synchronous feature-store round-trip on the critical path, and rebuilt around an async prefetch queue. p99 alert latency dropped from 4.2 seconds to 380 milliseconds with no additional hardware, and the on-call rotation stopped paging on transient batch backlog within the first week of the change shipping.

Vantage's engineering blog post on evaluating red-team coverage with class-balanced sampling matches a problem I hit directly: our own eval suite was systematically blind to a whole family of prompt-injection attacks because our sampling was frequency-weighted rather than class-weighted.

I would bring that same lesson, already paid for once, to your evaluation harness from day one.

I would welcome the chance to talk about where your detection roadmap is headed and where I could contribute first. I'm happy to walk through either project in more depth, or to bring a short write-up of the red-team harness to a first conversation.

Best regards,

Sarah Connor