

SARAH CONNOR

AI Security Engineer · Adversarial ML & Detection
sarah@example.com · Los Angeles, CA
example.com · github.com/sconnor · linkedin.com/in/sconnor

SUMMARY

AI security engineer specializing in adversarial ML detection and red-teaming for production LLM systems. Shipped detection infrastructure handling 50k+ req/s with sub-second alert latency.

EXPERIENCE

Senior AI Security Engineer Jun 2023 – Present
Cyberdyne Systems Los Angeles, CA

Cut adversarial-example detection latency 4x by rewriting the [batch, heads, seq, seq] scoring path to run inline with inference rather than as a post-hoc batch job.

- Built a red-team harness that found 12 prompt-injection classes missed by the prior eval suite, raising the pre-release catch rate from 61% to 94%.
Custom mutation-based fuzzer seeded from a 12k-prompt corpus, wired into the pre-release CI gate; findings triaged weekly with the app-sec team.
 - Led migration of the model-serving detection pipeline to a streaming architecture, cutting p99 alert latency from 4.2s to 380ms.
 - Mentored two junior engineers on adversarial evaluation methodology; both now own detection surfaces independently.
- PyTorch · vLLM · Triton · Kubernetes

AI Security Engineer Jan 2021 – May 2023
Widget Robotics Remote

Owned anomaly scoring for a fleet-telemetry pipeline processing 50k events/s.

- Shipped an anomaly scoring service that cut false-positive alert volume 70% by replacing static thresholds with a learned per-device baseline.
 - Built the eval harness now used across three teams to benchmark detection recall against a held-out attack corpus before every model release.
- Python · scikit-learn · Kafka · Terraform

Machine Learning Engineer Aug 2019 – Dec 2020
Nova Systems Analytics San Diego, CA

- Trained and deployed a fraud-scoring model that reduced chargeback losses 18% in its first quarter in production.
- Instrumented model-serving latency tracing, cutting p99 debugging time from days to hours for the on-call rotation.

SKILLS

Detection	Adversarial example detection, Anomaly scoring, Red-teaming, Eval design
Adversarial ML	Prompt injection, Model extraction, Data poisoning
Infra	Kubernetes, Terraform, Triton, CUDA, Kafka

EDUCATION

MS Computer Science 2019
University of California San Diego, US

CERTIFICATIONS

OSCP

Offensive Security Certified Professional

2022