

SARAH CONNOR

AI Security Engineer · Adversarial ML & Detection
sarah@example.com · Los Angeles, CA
example.com · github.com/sconnor · linkedin.com/in/sconnor

SUMMARY

AI security engineer specializing in adversarial ML detection and red-teaming for production LLM systems. Shipped detection infrastructure handling 50k+ req/s with sub-second alert latency.

EXPERIENCE

Senior AI Security Engineer
Cyberdyne Systems

Jun 2023 – Present
Los Angeles, CA

Cut adversarial-example detection latency 4x by rewriting the [batch, heads, seq, seq] scoring path to run inline with inference rather than as a post-hoc batch job.

- Built a red-team harness that found 12 prompt-injection classes missed by the prior eval suite, raising the pre-release catch rate from 61% to 94%.
Custom mutation-based fuzzer seeded from a 12k-prompt corpus, wired into the pre-release CI gate; findings triaged weekly with the app-sec team.
- Led migration of the model-serving detection pipeline to a streaming architecture, cutting p99 alert latency from 4.2s to 380ms.
- Mentored two junior engineers on adversarial evaluation methodology; both now own detection surfaces independently.

PyTorch · vLLM · Triton · Kubernetes

AI Security Engineer
Widget Robotics

Jan 2021 – May 2023
Remote

Owned anomaly scoring for a fleet-telemetry pipeline processing 50k events/s.

- Shipped an anomaly scoring service that cut false-positive alert volume 70% by replacing static thresholds with a learned per-device baseline.
- Built the eval harness now used across three teams to benchmark detection recall against a held-out attack corpus before every model release.

Python · scikit-learn · Kafka · Terraform

Machine Learning Engineer
Nova Systems Analytics

Aug 2019 – Dec 2020
San Diego, CA

- Trained and deployed a fraud-scoring model that reduced chargeback losses 18% in its first quarter in production.
- Instrumented model-serving latency tracing, cutting p99 debugging time from days to hours for the on-call rotation.

SKILLS

Detection	Adversarial example detection, Anomaly scoring, Red-teaming, Eval design
Adversarial ML	Prompt injection, Model extraction, Data poisoning
Infra	Kubernetes, Terraform, Triton, CUDA, Kafka

EDUCATION

MS Computer Science
University of California

2019
San Diego, US

> **CERTIFICATIONS**

OSCP

Offensive Security Certified Professional

2022