

# SARAH CONNOR

AI Security Engineer · Adversarial ML & Red Teaming  
sarah.connor@ena.one · +1 555 0142 · Austin, TX  
ena.one · [github.com/sconnor](https://github.com/sconnor) · [linkedin.com/in/sarahconnor](https://linkedin.com/in/sarahconnor)

---

## > SUMMARY

Security engineer specializing in adversarial robustness and red-teaming of production ML/LLM systems. 8+ years across offensive security and applied ML, with a focus on prompt-injection defense, model-extraction resistance, and secure MLOps pipelines.

## > EXPERIENCE

### Senior AI Security Engineer

Cyberdyne Systems

Jun 2023 – Present

Austin, TX

Leading red-team engagements against internal LLM-powered products before launch.

- Built an automated prompt-injection fuzzing harness that surfaced 40+ high-severity jailbreaks across 6 production LLM features prior to release.
- Led adversarial robustness reviews for a fraud-detection model serving 2M inference requests/day, reducing evasion success rate from 18% to under 2%.
- Designed and shipped a model-extraction detection pipeline (query-pattern anomaly scoring) now running across all customer-facing inference endpoints.

Python · PyTorch · Burp Suite · Kubernetes · OPA/Gatekeeper

### Security Engineer, Applied ML

Skynet Analytics

Aug 2020 – May 2023

Remote

Owned the security review process for the ML platform team's model deployment pipeline.

- Authored the org's first threat model for LLM-integrated services, adopted as the standard review template company-wide.
- Cut mean time to patch critical model-serving CVEs from 21 days to 4 by embedding security gating directly into the MLOps CI/CD pipeline.
- Ran quarterly red-team exercises against internal chatbots, discovering data-exfiltration paths later closed via output filtering and scoped tool permissions.

Terraform · AWS · MLflow · Semgrep · Snyc

### Penetration Tester

Omni Consumer Security

Jul 2017 – Jul 2020

Detroit, MI

Generalist offensive security consultant across web, network, and cloud engagements.

- Delivered 60+ client penetration tests across web apps, APIs, and AWS/Azure environments; consistently ranked top performer for report quality.
- Built an internal Burp Suite extension for automated JWT-tampering tests, adopted team-wide and cutting manual testing time by 30%.

Burp Suite · Nmap · Metasploit · AWS · Azure

## > PROJECTS

### promptfirewall

Open source

[github.com/sconnor/promptfirewall](https://github.com/sconnor/promptfirewall)

Model-agnostic prompt-injection classifier (400+ GitHub stars), used as a pre-filter in three published LLM security tools.

## > SKILLS

|          |                                                                         |
|----------|-------------------------------------------------------------------------|
| Security | Adversarial ML · Red teaming · Threat modeling · Penetration testing    |
| ML/AI    | LLMs · PyTorch · Prompt injection defense · Model extraction resistance |
| Infra    | Kubernetes · Terraform · AWS · CI/CD security gating                    |
| Practice | Secure MLOps · Vulnerability research · Incident response               |

## > EDUCATION

### M.S. Computer Science, Security concentration

University of Texas at Austin

2017

Austin, TX

**B.S. Computer Engineering**  
*Georgia Institute of Technology*

2015  
*Atlanta, GA*

>

## **CERTIFICATIONS**

---

**OSCP (Offensive Security Certified Professional)**  
*Offensive Security*

2019  
*Online*

**GIAC Certified Incident Handler (GCIH)**  
*GIAC*

2021  
*Online*