

SARAH CONNOR ■

AI Security Engineer · Adversarial ML & Red Teaming
sarah.connor@ena.one · +1 555 0142 · Austin, TX
ena.one · github.com/sconnor · linkedin.com/in/sarahconnor

> SUMMARY

Security engineer specializing in adversarial robustness and red-teaming of production ML/LLM systems. 8+ years across offensive security and applied ML, with a focus on prompt-injection defense, model-extraction resistance, and secure MLOps pipelines.

> EXPERIENCE

Senior AI Security Engineer

Cyberdyne Systems

Jun 2023 – Present

Austin, TX

Leading red-team engagements against internal LLM-powered products before launch.

- Built an automated prompt-injection fuzzing harness that surfaced 40+ high-severity jailbreaks across 6 production LLM features prior to release.
- Led adversarial robustness reviews for a fraud-detection model serving 2M inference requests/day, reducing evasion success rate from 18% to under 2%.
- Designed and shipped a model-extraction detection pipeline (query-pattern anomaly scoring) now running across all customer-facing inference endpoints.

Python · PyTorch · Burp Suite · Kubernetes · OPA/Gatekeeper

Security Engineer, Applied ML

Skynet Analytics

Aug 2020 – May 2023

Remote

Owned the security review process for the ML platform team's model deployment pipeline.

- Authored the org's first threat model for LLM-integrated services, adopted as the standard review template company-wide.
- Cut mean time to patch critical model-serving CVEs from 21 days to 4 by embedding security gating directly into the MLOps CI/CD pipeline.
- Ran quarterly red-team exercises against internal chatbots, discovering data-exfiltration paths later closed via output filtering and scoped tool permissions.

Terraform · AWS · MLflow · Semgrep · Snyk

Penetration Tester

Omni Consumer Security

Jul 2017 – Jul 2020

Detroit, MI

Generalist offensive security consultant across web, network, and cloud engagements.

- Delivered 60+ client penetration tests across web apps, APIs, and AWS/Azure environments; consistently ranked top performer for report quality.
- Built an internal Burp Suite extension for automated JWT-tampering tests, adopted team-wide and cutting manual testing time by 30%.

Burp Suite · Nmap · Metasploit · AWS · Azure

> PROJECTS

promptfirewall

Open source

github.com/sconnor/promptfirewall

Model-agnostic prompt-injection classifier (400+ GitHub stars), used as a pre-filter in three published LLM security tools.

> SKILLS

Security	Adversarial ML · Red teaming · Threat modeling · Penetration testing
ML/AI	LLMs · PyTorch · Prompt injection defense · Model extraction resistance
Infra	Kubernetes · Terraform · AWS · CI/CD security gating
Practice	Secure MLOps · Vulnerability research · Incident response

> EDUCATION

M.S. Computer Science, Security concentration

University of Texas at Austin

2017

Austin, TX

B.S. Computer Engineering
Georgia Institute of Technology

2015
Atlanta, GA

>

CERTIFICATIONS

OSCP (Offensive Security Certified Professional)
Offensive Security

2019
Online

GIAC Certified Incident Handler (GCIH)
GIAC

2021
Online