

SARAH CONNOR ■

AI Security Engineer · Adversarial ML & Red Teaming

sarah.connor@ena.one · +1 555 0142 · Austin, TX

ena.one · github.com/sconnor · linkedin.com/in/sarahconnor

> 28 August 2026

Dear Hiring Manager,

Cortex Labs' recent engineering post on the automated red-teaming pipeline behind your assistant launch is the reason I'm writing today rather than just bookmarking your careers page. Adversarial robustness for production LLM systems is the specific problem I've spent the last three years on, and it is rare to see a team publish the actual mutation strategy behind their fuzzing harness instead of a marketing summary of "safety testing."

At Cyberdyne Systems I built and ran exactly that kind of system: an automated prompt-injection fuzzing harness driven by a custom mutation engine over a 12,000-seed corpus, wired directly into the release-gate CI job rather than run as a separate, easily-skipped audit step. The harder problem wasn't finding jailbreaks — the harness surfaced 40+ high-severity ones across six production features in its first quarter — it was building a triage process the app-sec team could sustain weekly without burning out on false positives, which meant investing as much in the scoring and deduplication logic as in the mutations themselves.

Your team's public work on constrained decoding for tool-call outputs tackles the same class of problem from the model side rather than the harness side, and I'd like to compare notes on where the two approaches overlap. The scoped-tool-permission pattern in your last release notes also mirrors a mitigation I shipped after a red-team engagement uncovered a similar data-exfiltration path in an internal chatbot.

I'd welcome the chance to talk through how this experience could plug into Cortex Labs' red-teaming roadmap. I'm available for a call any weekday and can share the fuzzing harness's design doc on request.

Best regards,
Sarah Connor